

Pacman Capture the Flag Assignment 4 Report

GROUP 03

Shekhar Devm Upadhyay	Aiman Shenawa
20010531	001021
sdup@kth.se	ashenawa@kth.se



May 2023

Abstract

This report presents a study on the development of AI agents for controlling a team of two in a competitive Pac-Man Capture the Flag (CTF) scenario. The objective of the agents is to maximize their score by efficiently collecting food from the enemy turf within a limited number of moves. This report studies various approaches used in the field of artificial intelligence, including minimax game tree search, modifications like alpha-beta pruning and move ordering, Monte Carlo Tree Search (MCTS), Expectimax, and decision-tree-based approaches, presents the final solution employed by the authors, and does a comparative analysis of the approaches utilized by all the teams in the competition. By analyzing and comparing these strategies, we aim to identify effective techniques for AI-controlled agents in Pac-Man CTF.

1 Introduction

With the continuous development of society and technology, the need for autonomous agents is increasing. It is becoming more and more important to develop agents that can operate in dynamic and complex environments autonomously. In this report we will explore the problem of developing an agent that can play the game of Pacman capture the flag. This is a project developed by UC Berkely, it is a challenge that combines classic video game mechanics with artificial intelligence techniques. The goal of the project is to develop an agent that can play the game of Pacman capture the flag. The game is played on a grid, where the agent can move in four directions, up, down, left and right. The grid is split into two halves, in each half there is a team of Pacman agents, the goal of the game is to capture the food pellets from the other team's side and bring them back to your side while also preventing the other team from doing the same. Multiple aspects of artificial intelligence will be explored in this report, such as minimax game tree search, Monte Carlo Tree Search (MCTS), Expectimax, and decision-tree-based approaches. The final solution included the use of Minimax and alpha-beta pruning.



Figure 1: Illustration of environment

The figure shows the game environment, the blue and red dots represent the food pellets, the blue and red lines represent the walls. As seen in the figure above there are both Pacman and ghosts. The agent is a ghost when defending its own side. When crossing over to the enemy side to collect pellets it becomes a pacman. An agent can go into enemy territory and collect pellets, but the score will only increase when the agent returns back to its own side. In the figure there are also white pellets, those are power pellets, when an agent collects a power pellet it can eat the enemy agents for a period of 40 moves. During that period the enemy ghost agents are regarded as scared. When an agent is eaten it respawns at its own side, if the agent was scared before getting eaten it will no longer be scared after respawning.

The information each agent has access to is the following:

- The friend agent's positions
- The total score
- The position of all the food pellets on the map, this is updated every time an agent moves
- The position of all the power pellets on the map, this is updated every time an agent moves
- The position of an enemy agent if it is within 5 squares of the agent or a teammate

1.1 Task Description

This report will discuss the problem of developing an agent that can play the game of Pacman capture the flag. The aim is to develop two agents which are able to defend their side of the map and also collect food pellets from the enemy side. The agents will be able to move in four directions, up, down, left and right. The agents

These problems and their solutions are addressed in further detail in the following section 2.

1.2 Contribution

In this report we will implement Minimax and alpha-beta pruning into the UC Berkeley Pacman Capture the Flag project. We will also explore the use of Monte Carlo Tree Search (MCTS), Expectimax, and decision-tree-based approaches. We will also compare the results of our solution with the results of other methods used for the same problem.

An essential aspect of Pacman capture the flag is the ability to cooperate as a team and communicate with the other agent. Potential real world applications of this project are many, as the ability of intelligent agents to make strategic decisions in a dynamic environment has many applications. Applications include multi-robot system coordination, autonomous vehicles, and unmanned aerial vehicles.

1.3 Outline

This report is organized as follows. First we introduce the problem and its importance in section ???. Then we discuss and highlight the related work and required background knowledge in order to understand the proposed solution in section 2. This is followed by a detailed description of the proposed solution in section 3. Next we present the experimental setup and the results in section 4. Finally, we summarize the results and conclude the report in section 5.

2 Related Work

The success of AI agents in Pac-Man CTF heavily relies on the choice and implementation of appropriate strategies. This section provides an overview of the relevant research in this domain.

2.1 Minimax Game Tree Search

The problem of finding the optimal move in a game can be formulated as a search problem. The search space is a tree, where each node represents a state of the game, and each edge represents a possible move. The root node represents the current state of the game, and the leaf nodes represent the terminal states of the game. The goal is to find the optimal path from the root node to a leaf node.

The minimax algorithm (shown in Algorithm 1) is a recursive algorithm that computes the optimal move for a player in a two-player game. The algorithm assumes that the opponent plays optimally, and it tries to minimize the maximum loss that can be incurred by the player. It is a classic approach in Game Theory, aiming to find the optimal move in a two-player, zero-sum game.

By constructing a game tree and evaluating the utility of different game states, the algorithm enables agents to make informed decisions.

Sometimes, however, it is not feasible to search the entire game tree, due to the large number of possible moves. To address this issue, the algorithm can be modified to limit the depth of the search tree. The algorithm can also be modified to incorporate heuristics that evaluate the utility of non-terminal game states. Variations of minimax, such as alpha-beta pruning, move ordering based on heuristics, iterative deepening search (IDS), and Repeated States Checking (RSC), have been proposed to improve its efficiency and effectiveness [1][2][3].

Algorithm 1 Minimax Algorithm

```

1: procedure MINIMAX(node, depth, maximizingPlayer)
2:   if depth = 0 or node is a terminal node then
3:     return the heuristic value of node
4:   end if
5:   if maximizingPlayer then
6:     value  $\leftarrow -\infty$ 
7:     for each child of node do
8:       value  $\leftarrow \max(\text{value}, \text{Minimax}(\text{child}, \text{depth} - 1, \text{False}))$ 
9:     end for
10:    return value
11:  else
12:    value  $\leftarrow \infty$ 
13:    for each child of node do
14:      value  $\leftarrow \min(\text{value}, \text{Minimax}(\text{child}, \text{depth} - 1, \text{True}))$ 
15:    end for
16:    return value
17:  end if
18: end procedure

```

2.2 Expectimax

Expectimax is a variant of the minimax algorithm that considers uncertain outcomes in games. It is commonly used in domains with probabilistic elements. In Pac-Man CTF, Expectimax can be employed to account for the ghost movement and uncertain states, enabling agents to make rational decisions under uncertainty [6].

2.3 Monte Carlo Tree Search

MCTS is a sampling-based search algorithm that has demonstrated remarkable success in various game-playing domains. By combining tree exploration with random rollouts, MCTS performs effective exploration and exploitation of the search space. This approach has been applied to Pac-Man CTF, where agents simulate multiple playouts to estimate the value of different actions and make informed decisions [4][5].

2.4 Decision Tree Based Approaches

Decision trees offer a structured representation of decision-making processes. By learning from historical data and constructing decision rules, decision-tree-based approaches provide an interpretable framework for agent behavior. These techniques have been explored in Pac-Man CTF, where agents use decision trees to guide their actions based on various game state features [7][8].

2.5 Hybrid Approaches

Several studies have proposed hybrid strategies by combining different AI techniques in Pac-Man CTF. For example, integrating MCTS with minimax or Expectimax algorithms has shown promise in improving agent performance [9][10]. Such hybrid approaches leverage the strengths of different methods to create robust and adaptable AI agents.

3 Method

In this section, we will discuss ...

4 Experiments and Results

First, we will describe the experimental setup for the two problems, and then we will present the results for each of the problems separately.

4.1 Experimental Setup

The experimental environment was largely dependent on the problems. Both problems shared a similar map and obstacle structure, the main difference being the agent objectives, environment sizes, and the fact that we didn't have any obstacles in the soccer environment (except, of course, for the boundaries of the field). For each of the problems, we were to compute and pass a set of inputs to the agents (different based on whether the agent was a car or a drone). The agents moved based on the provided input, and the updated positions and velocities were used in the next timestep to compute the next set of inputs. The agents were also provided with a set of parameters that decided how the different forces were weighted. The parameters were tuned by hand separately for each type of agent, and were not changed during the course of the simulation.

All the solutions for Collision Avoidance were repeatedly tested in the Unity game environment on three terrains: an open field, an intersection, and a random map. The open field was a simple map with no obstacles. The intersection was a map with a crossroad and a few obstacles. The random map was a map with randomly generated obstacles. There were two possible initializations for the positions of the agents: first, in a circular formation, and second, in a random formation. The solutions were tested in the Unity game environment with 50 agents. Based on performance, the solutions were improved and adjusted in an iterative manner.

4.2 Results

In this section, we summarize the results of our experiments. We first present the results of the Collision Avoidance problem, and then present the results of the Soccer problem. We then go on to discuss the merits and shortcomings of the best solutions from all the groups.

5 Summary and Conclusions

A Appendix